

A forecast is a range.

A stock tip is one number. What happens next is a range. We forecast that range for 60 NSE stocks, 21 trading days ahead.



5 min

Executive pitch


10 min

Technical deep dive

Live

Demo

Advisers get a tip. They need the range, the reason and the size.



- What they have today
- Tips and single-number targets
 - No view of the range of outcomes
 - No link to the client's risk tolerance

467 adviser hours back each year: a 483% return on cost.

Stated assumption

200 clients × 4 reviews a year

× 45 down to 10 minutes each

× Rs 1,500 an adviser hour

against Rs 1.2 lakh running cost a year

What follows

800
reviews a year

467
adviser hours saved

Rs 7.0 lakh
of adviser time a year

483%
return on cost

Route to market

Free Research tool for self-directed investors

Paid seat Advisers who need client-ready evidence sheets

Channel Brokers, through the Upstox and Zerodha quote adapters already built

Not claimed: extra return

On unseen data the top-5 picks returned −4.4% a year against +3.3% for holding every stock. We sell speed, discipline and honest ranges.

From raw prices to the app, with leakage designed out.



Data

60

liquid NSE large caps, five in each of 12 sectors

118,641


stock-days of daily prices, Feb 2018 to
Oct 2026

21 days

ahead: the return the model is asked to forecast

Feature selection

score = relevance - 0.5 × redundancy




10 of 19 kept. RSI, stochastic %K and Williams %R track the same thing, so one stays.

Fundamentals left out of the model. Only today's ratios exist, so using them on past dates would leak the future.

Tuning beat the notebook by 9.2%, and only tied the naive baseline.

Purged walk-forward cross-validation: five expanding folds



Cross-validated pinball loss lower is better

Naive baseline		0.02395
Ridge	Point forecast only, so no quantile loss. MAE 0.0731, naive 0.0722.	
LightGBM, notebook settings		0.02648
LightGBM, tuned	chosen	0.02406

375

model fits: 25 settings × 5 folds × 3 quantiles

-9.2%

pinball loss against the notebook settings


+0.5%

against naive: a tie, inside fold-to-fold variation

The search chose 4 leaves and 100 trees. With a weak signal, bigger trees fit noise.

Honest ranges. No market-beating picks.

Unseen data: outcomes inside the P10 to P90 range



0.01880

pinball loss, naive
0.01915


+0.046

rank IC, positive on 56%
of 407 days

51.0%

direction right

Backtest on the same unseen period



-4.4% a year

Top-5 picks


+3.3% a year

All 60 stocks, equal weight

Median of 21 start days, after costs.
Picks won from 24% of them.


What drives a forecast (SHAP)

A higher value pushes it ← down | up →



Where it is wrong

Rank IC by market regime, unseen data



Direction right **31.1%** when the market fell,
64.6% when it rose.

Limits

- 21-day returns are mostly noise
- 60 of today's large caps: survivorship bias
- No taxes or slippage beyond a flat cost

Next: point-in-time fundamentals, a larger universe, paper trading.

Live demo

Profile intake

Ranked book

Stock research

Portfolio lab

Model report



risk and capital

signal and range

SHAP for one stock

position sizes

the numbers shown here

Live app

[**fn.iimbg.com**](https://fn.iimbg.com)

Landing page

[**fin.iimbg.com**](https://fin.iimbg.com)

Code

[**github.com/build-with-bala/ai-portfolio-advisor**](https://github.com/build-with-bala/ai-portfolio-advisor)

22

automated tests

1 command

retrains the model in about a minute

Built with Claude Code

prompts, failures and fixes are in docs/VIBE_CODING_LOG.md. The team owns and defends the method.